

# Parallel Programming Models

Lecture 3  
Introduction to High Performance Computing  
IN4049 TUDelft  
2013

# Recap Lecture 2 [1]

---

*Introduction to HPC  
in4049 TU Delft  
2013*

- Design parallel applications
  - *Systematic process*
  - Performance evaluation using Amdhal's law
    - *First estimates are available at design time!*
  - Models of parallel architectures
    - *Shared memory*
    - *Distributed memory*
    - *Virtual shared memory*

## Models of parallel computation

- Conceptual level: closer to the application
  - *farmer/worker*
  - *divide and conquer*
  - *data parallelism*
  - *task parallelism*
  - *bulk synchronous*
- System level: closer to the machine
  - *(Logical) data spaces and programs*
  - *Communication and synchronization*

# Today – Lecture 3

Introduction to HPC  
in4049 TU Delft  
2013

- Parallel programming models
  - Classic: OpenMP, HPF, MPI
  - Low-level: pthreads
  - Modern reincarnations: TBB, ArBB, Cilk (Satin), Titanium, UPC, MapReduce (lecture 4)
  - HW-centric: CUDA, OpenCL, OpenACC, ... (lecture 4)
- Reading assignment (optional): on comparing performance models:
  - *Varbanescu et. al - “Towards an Effective Unified Programming Model for Many-Cores”, IPDPS Workshops 2012*

---

# Parallel programming models

---

# Parallel Programming Model

Introduction to HPC  
in4049 TU Delft  
2013

- Programming model = language + libraries that create a model of computation (i.e., abstract machine)
- Abstraction level = qualifier for the distance between the abstract and the real machine
  - *E.g. Assembly language: very low vs. Java: high/very high*
  - *More difficult to estimate for parallel models*
- Empirical evaluation:
  - *High-level languages:*
    - *Pro: Programmer friendly*
    - *Con: Performance impacted by overheads*
  - *Low-level languages:*
    - *Pro: Performance*
    - *Con: Difficult to learn and use*

- Control model
  - Parallel operations creation
  - (Logical) Ordering
- Data model
  - Shared vs. Private data
  - Logical access to data (i.e., communication)
- Synchronization model
  - Coordinating parallelism
- Cost model(/performance impact)
  - Expensive vs. Cheap operations
  - Overhead

---

# Classical models: OpenMP

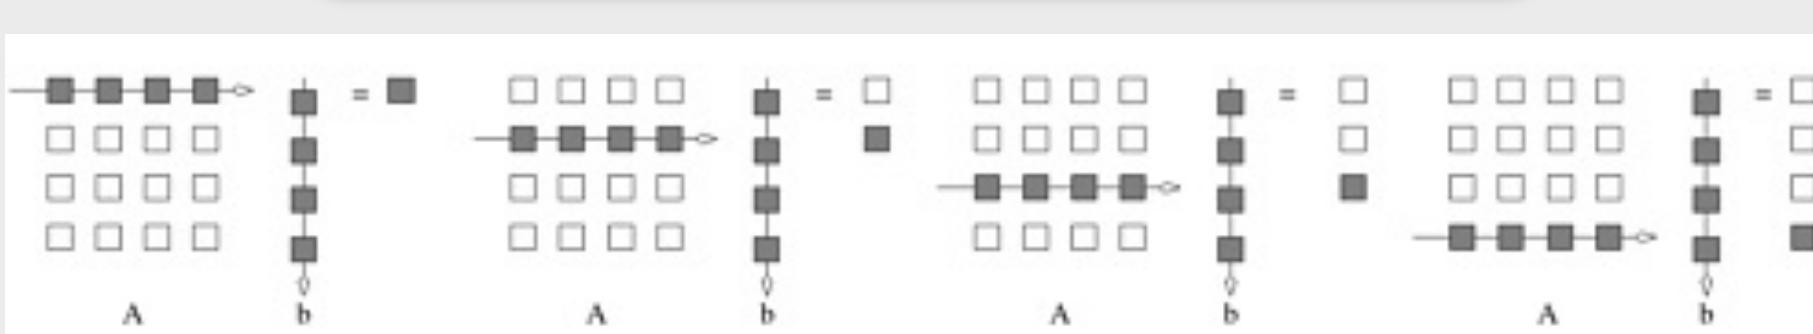
---

# App: Matrix-Vector Product (MVP)

Introduction to HPC  
in4049 TU Delft  
2013

- On the forthcoming slides this operation is used for demonstration of the different parallel languages

```
for(i=0; i<dim; i++) {  
    y[i]=0;  
    for(j=0; j<dim; j++) {  
  
        y[i]=y[i] + A[i,j] * x[j];  
    }  
}
```



- Control:
  - *Parallel Loops*
  - *Parallel Sections*
- Data:
  - *Shared and/or private, explicitly specified*
- Synchronization
  - *Explicit, using critical sections*
  - *Atomic operations*
- Communication
  - *Shared variables*

- Extension to C, C++, Fortran (pragma-based)
- General format: #pragma omp directive [clauses]
- Number of threads: compile-time / run-time
- Example\*

```
1. int main (int argc, char *argv[]) {  
2.     int nthreads, tid;  
3.     #pragma omp parallel private(nthreads, tid) {  
4.         tid = omp_get_thread_num();  
5.         printf("Hello World from %d\n", tid);  
6.         nthreads = omp_get_num_threads();  
7.         if (tid == 0)  
8.             printf("Threads = %d\n", nthreads);  
9.     }
```

\*More: <https://computing.llnl.gov/tutorials/openMP/exercise.html>

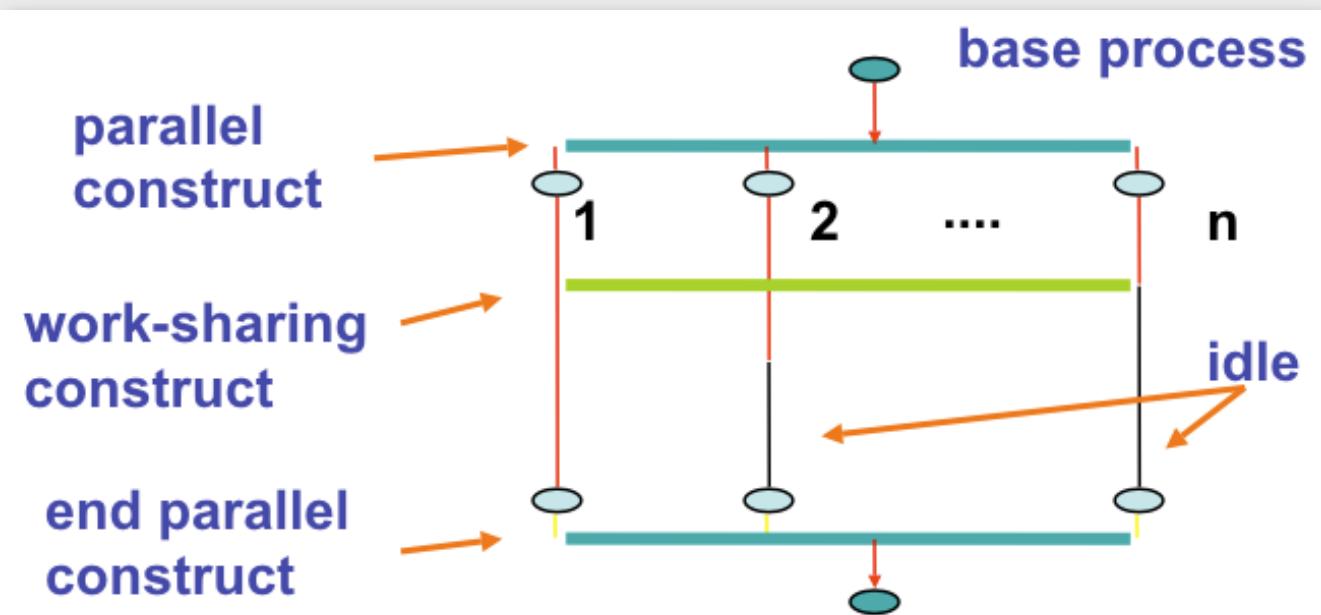
# OpenMP execution model

Introduction to HPC  
in4049 TU Delft  
2013

Parallel construct creates a number of worker processes (or threads)

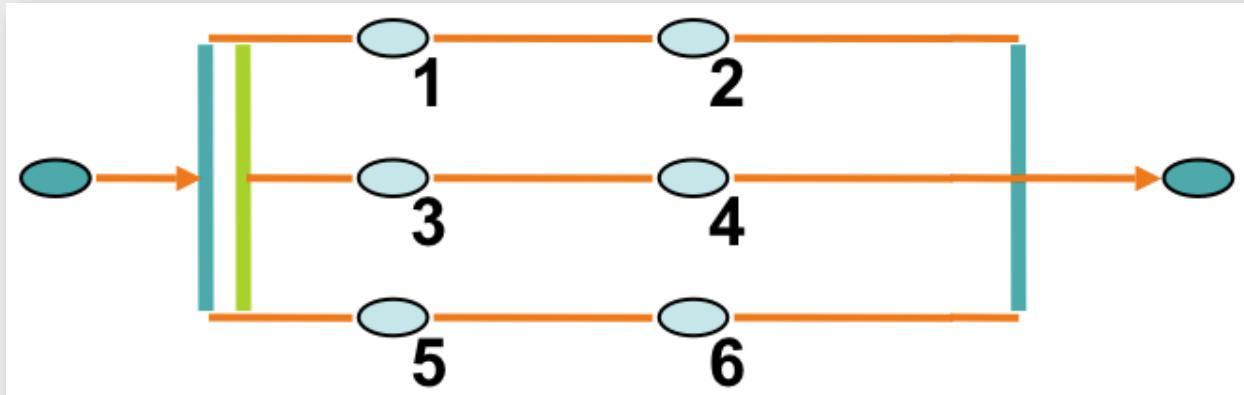
Work-sharing construct hands out work to available workers (threads)

- #pragma omp sections [clauses]
- #pragma omp for [clauses]



# OpenMP: parallel for

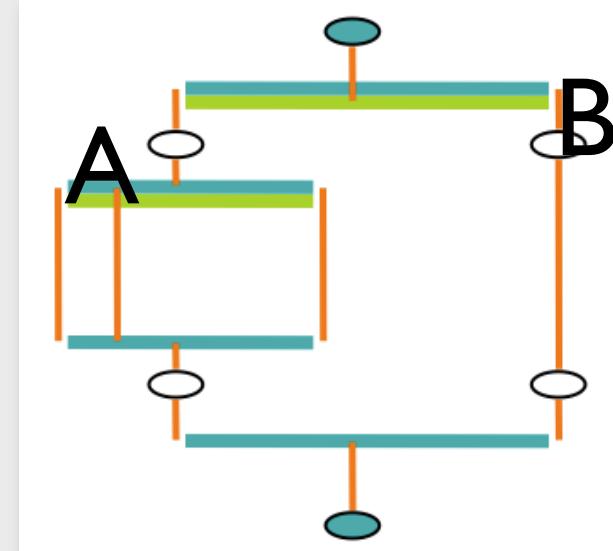
```
#pragma omp parallel
    #pragma omp for schedule(static,2)
        for(i=0;i<dim;i++) { ... }
```



# OpenMP: parallel section

- Combines parallel and work-sharing construct
- Allows independent execution of subprograms
  - *Task parallelism*

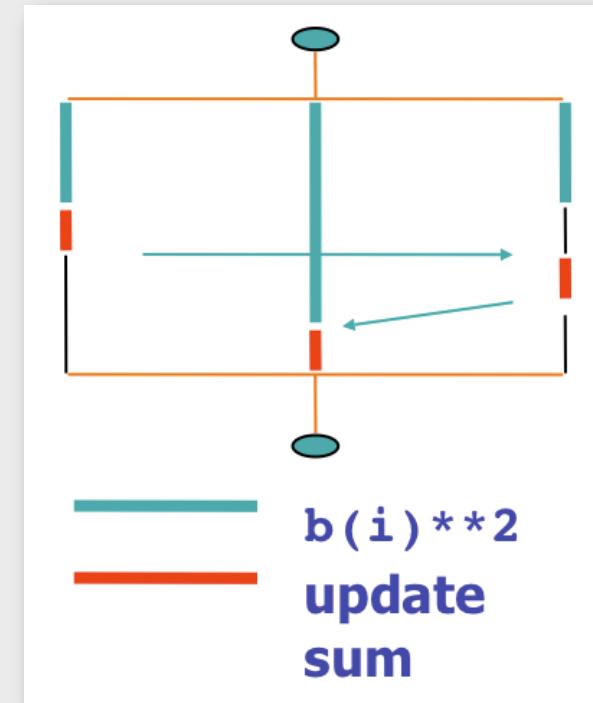
```
#pragma omp parallel sections {
    #pragma omp section {
        task_A();           // parallel
    for
        }
    #pragma omp section  {
        task_B();           // simple
    }
}
```



# OpenMP: locks and critical sections

Introduction to HPC  
in4049 TU Delft  
2013

```
for(i=0;i<dim;i++) {  
    sum = sum + b[i]**2;  
}
```



```
#pragma omp parallel private(i) reduction(+:sum)  
schedule(dynamic,1) {  
    for(i=0;i<dim;i++) { sum = sum + b[i]**2 }  
}
```

# OpenMP: MVP example

Introduction to HPC  
in4049 TU Delft  
2013

```
for(i=0; i<dim; i++) {  
    y[i]=0;  
    for(j=0; j<dim; j++) {  
        y[i]=y[i] + A[i,j] * x[j];  
    }  
}
```

```
#pragma omp parallel for private(j)  
for(i=0; i<dim; i++) {  
    y[i]=0;  
    for(j=0; j<dim; j++) {  
        y[i]=y[i] + A[i,j] * x[j];  
    }  
}
```

# OpenMP: MVP example

Introduction to HPC  
in4049 TU Delft  
2013

```
for(i=0; i<dim; i++) {  
    y[i]=0;  
    for(j=0; j<dim; j++) {  
        y[i]=y[i] + A[i,j] * x[j];  
    }  
}
```

```
for(i=0; i<dim; i++) {  
    tmp=0;  
    #pragma omp parallel for reduction(+ :tmp)  
    for(j=0; j<dim; j++) {  
        tmp=tmp + A[i,j] * x[j];  
    }  
    y[i] = tmp;  
}
```

- Shared memory programming model
- Explicit parallelism
  - Fork-join model
- Implicit data distribution
- Allows for data and task parallelism
- Shared and private variables
  - Communication via shared variables
- Synchronization by critical sections
- High abstraction level
- Preserves sequential code (when possible/needed)

# OpenMP Quiz

Introduction to HPC  
in4049 TU Delft  
2013

Which of the following *conceptual programming models* be represented in OpenMP:

- *farmer/workers*
- *divide and conquer*
- *data parallelism*
- *task parallelism*
- *bulk synchronous*
- Write a short (pseudo-code) example for each one of those for which this is possible.
- Indicate the functionality.
- Syntax correctness is not an issue.

---

# Low-level model: pthreads

---

# Shared memory: pthreads

Introduction to HPC  
in4049 TU Delft  
2013

- Applications are divided into threads
- Each thread receives its own functionality
  - SPMD – Single process, multiple data

```
for (i=0; i<N; i++) {  
    create_thread(i, function[i]);  
    // run_thread(i)  
    join_thread(i); }
```

- MPMD – Multiple process, multiple data

```
create_thread(1, function_1);  
create_thread(2, function_2);  
  
join_thread(1); join_thread(2);
```

- Thread management is explicit
  - Creation, joining, killing
- Thread scheduling is (typically) done by OS
  - No pre-determined model
- Communication
  - Explicit via shared variables
  - Protection for race conditions needed
  - Protection against deadlocks
- Synchronization
  - Locks
  - Barriers (not available in the standard!)

# Pthreads Quiz

Introduction to HPC  
in4049 TU Delft  
2013

Are there any models that cannot be expressed in OpenMP,  
but can be expressed in pThreads?

- *farmer/workers*
- *divide and conquer*
- *data parallelism*
- *task parallelism*
- *bulk synchronous*

---

# Classical models: HPF

---

# (Logically) Distributed memory: HPF

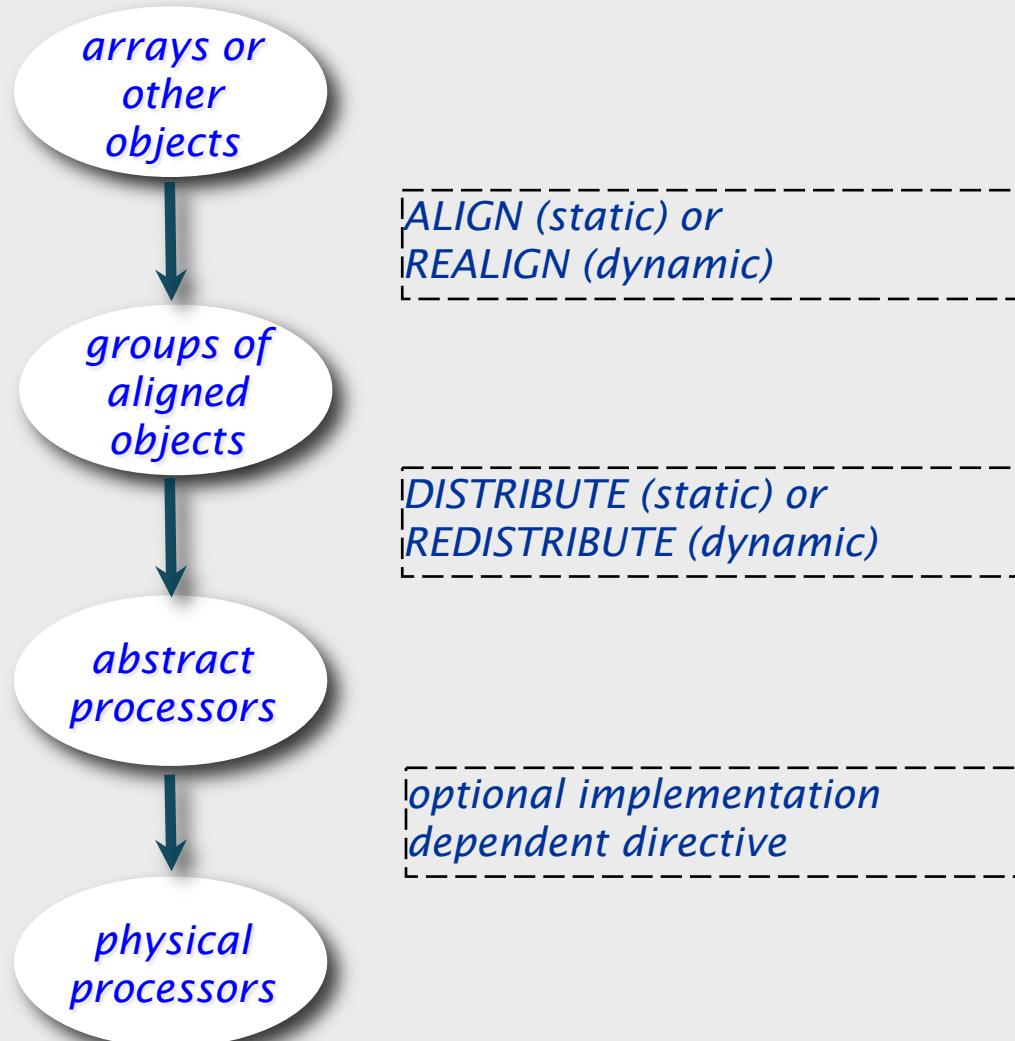
Introduction to HPC  
in4049 TU Delft  
2013

- High Performance Fortran
- Based on adding annotations to sequential Fortran90 base program
- Main features
  - *parallel constructs*
  - *work-sharing based on data distribution statements*
- Data parallel model
  - *Distributes the array elements over virtual processors*

- Control:
  - *Data-driven: computation follows data distribution*
  - *Parallel arrays and loops*
- Data:
  - *Distributed memory model*
  - *Full arrays still available for the applications, thus ...*
- Synchronization and communication
  - *Implicit*
  - *Generated whenever non-local data is accessed.*

# HPF: data distribution model

Introduction to HPC  
in4049 TU Delft  
2013



# HPF: data mapping example

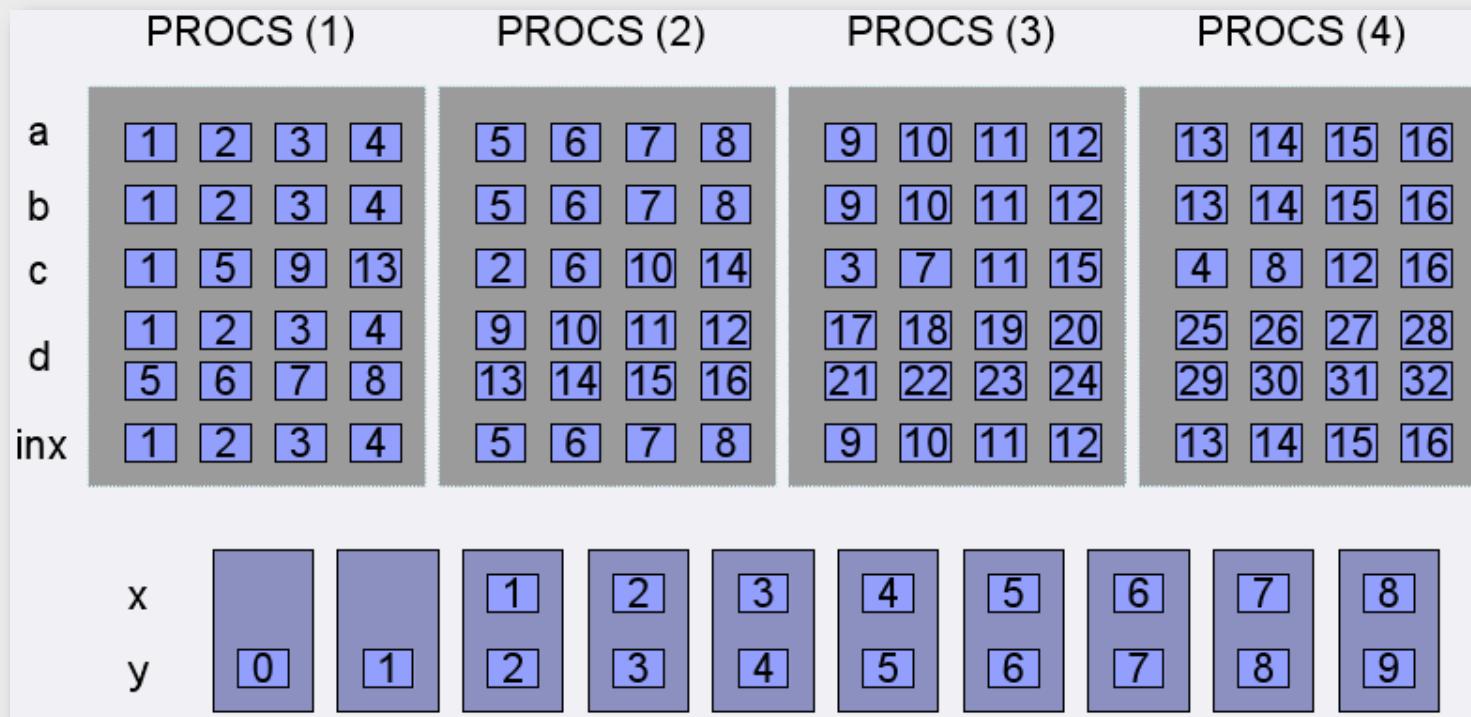
*Introduction to HPC  
in4049 TU Delft  
2013*

```
REAL, DIMENSION (16) :: A, B, C
REAL, DIMENSION (32) :: D
REAL, DIMENSION (8) :: X
REAL, DIMENSION (0:9) :: Y
INTEGER, DIMENSION (16) :: INX
!HPF$ PROCESSORS, DIMENSION(4) :: PROC
!HPF$ DISTRIBUTE, (BLOCK) ONTO PROCS :: A, B, D, INX
!HPF$ DISTRIBUTE, (CYCLIC) ONTO PROCS :: C
!HPF$ ALIGN X(I) WITH Y(I+1)
```

# HPF: data mapping example

Introduction to HPC  
in4049 TU Delft  
2013

- the code in the previous slide is resulting as shown here:

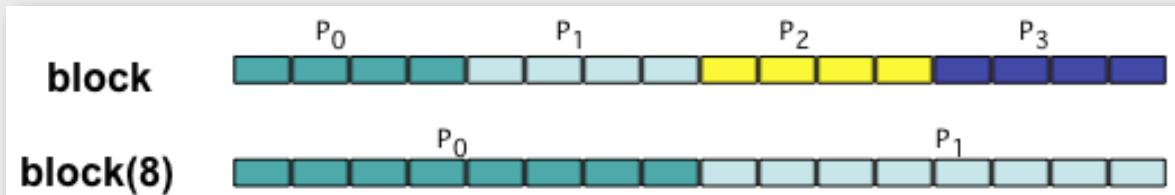


- Distributes the array elements over virtual processors
- Various predefined distribution functions
  - *block, block(m)*
  - *cyclic, cyclic(m)*
  - *\** (*all elements in dimension to same processor*)
- examples:

```
!HPF$ PROCESSORS, DIMENSION(4,4) ONTO p  
  
!HPF$ PROCESSORS, DIMENSION(16) ONTO q  
  
!HPF$ DISTRIBUTE a(block,block) ONTO p  
  
!HPF$ DISTRIBUTE a(cyclic,*) ONTO q  
  
!HPF$ DISTRIBUTE a(*,*) ONTO powerpc
```

- BLOCK means give equal sized contiguous blocks of array elements to processors
- Suited for problems with equal computational load per element
- Example:

```
REAL a(16)
!HPF$      PROCESSORS p(4)
!HPF$      DISTRIBUTUTE a(block) ONTO p
or
!HPF$      DISTRIBUTUTE a(block(8)) ONTO
p
```



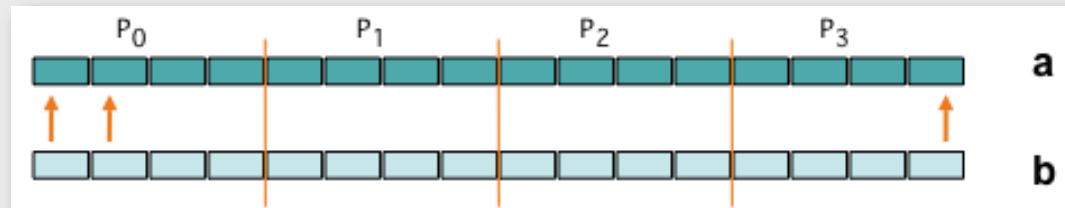
- Arrays can be ALIGNed to each other
- Aligned arrays are collectively distributed through top array
- Top array can be a TEMPLATE (no data, only shape)
- Purpose:
  - *Specify locality: two aligned array elements reside on the same processor*
  - *Specify replication and collapsing of dimensions*
  - *distribute a group of arrays with a single DISTRIBUTE statement*

# HPF alignment example (1)

Introduction to HPC  
in4049 TU Delft  
2013

- Example: one-to-one alignment

```
REAL a(16), b(16)
!HPF$ PROCESSORS p(4)
!HPF$      ALIGN b(:) WITH a(:)
!HPF$      DISTRIBUTE a(block) ONTO
p
```

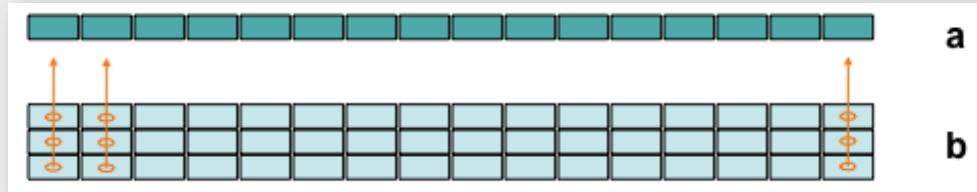


# HPF alignment example (2)

Introduction to HPC  
in4049 TU Delft  
2013

- Example: collapsed alignment

```
REAL a(16), b(16,3)
!HPF$ PROCESSORS p(4)
!HPF$      ALIGN b(:,*) WITH a(:)
!HPF$      DISTRIBUTE a(block) ONTO p
```



- F90 array assignment statements
- FORALL statement
- Examples:

*array statements*

```
REAL, DIMENSION(N)::a,b,c
a= b+c
a(1:10:2)=b(6:10) + c(1:5)
```

*FORALL statement*

```
FORALL (i=1:n, j=1:m, a(i,j).NE.0)
    a(i,j) = 1.0 / a(i,j)
```

# HPF example: MVP

Introduction to HPC  
in4049 TU Delft  
2013

```
REAL a(m,n), y(m), x(n)
!HPF$      PROCESSORS p(4)
!HPF$      ALIGN a(:,*) WITH y(:)
!HPF$      ALIGN x(:) WITH y(:)
!HPF$      DISTRIBUTE y(block) ONTO p

!HPF$ INDEPENDENT
DO j=1,n
    FORALL (i=1:m)
        y(i) =y(i) + a(i,j)*x(j)
    END FORALL
END DO
```

# HPF example: MVP

Introduction to HPC  
in4049 TU Delft  
2013

```
REAL a (m,n) , y (m) , x (n)
!HPF$      PROCESSORS p (4)
!HPF$      ALIGN a (:, *) WITH y (:)
!HPF$      ALIGN x (: ) WITH y (: )
!HPF$      DISTRIBUTE y (block) ONTO p

FORALL (i=1,m)
    y (i) = DOT_PRODUCT (a (i, :), x (: ))
```

- Data-parallel programming model
  - Computation follows data
  - Communication and synchronization “by request”
- Data distribution & alignment
  - Essential for performance, locality, low communication/synchronization
- Different processors view
  - Allow for easier data distributions
- High abstraction level
- Performance depends on compiler
  - Large overheads depending on the machine

## HPF Quiz:

---

Introduction to HPC  
in4049 TU Delft  
2013

Which of the following *conceptual programming models* be represented in HPF:

- *farmer/workers*
- *divide and conquer*
- *data parallelism*
- *task parallelism*
- *bulk synchronous*
- Write a short (pseudo-code) example for each one of those for which this is possible.
- Indicate the functionality
- Syntax correctness is not an issue.

---

# Classical models: Message Passing (MPI)

---

- Control:
  - *Independent processes*
  - *Typical SPMD*
    - *MPMD by process ID*
- Data:
  - *No notion of global data*
  - *Every process has local copies of the data*
- Communication
  - *By message passing*
  - *Variants: 1-1, 1-n, n-1, n-n*
- No special synchronization constructs
  - *implicit synchronization by message transfer*
  - *barrier synchronization is done by all-to-all broadcast*
- Abstraction level: low/medium

# Message Passing: MVP Example

Introduction to HPC  
in4049 TU Delft  
2013

```
for(i=0; i<dim; i++) {  
    y[i]=0;  
    for(j=0; j<dim; j++) {  
        y[i]=y[i] + A[i,j] * x[j]; } }
```

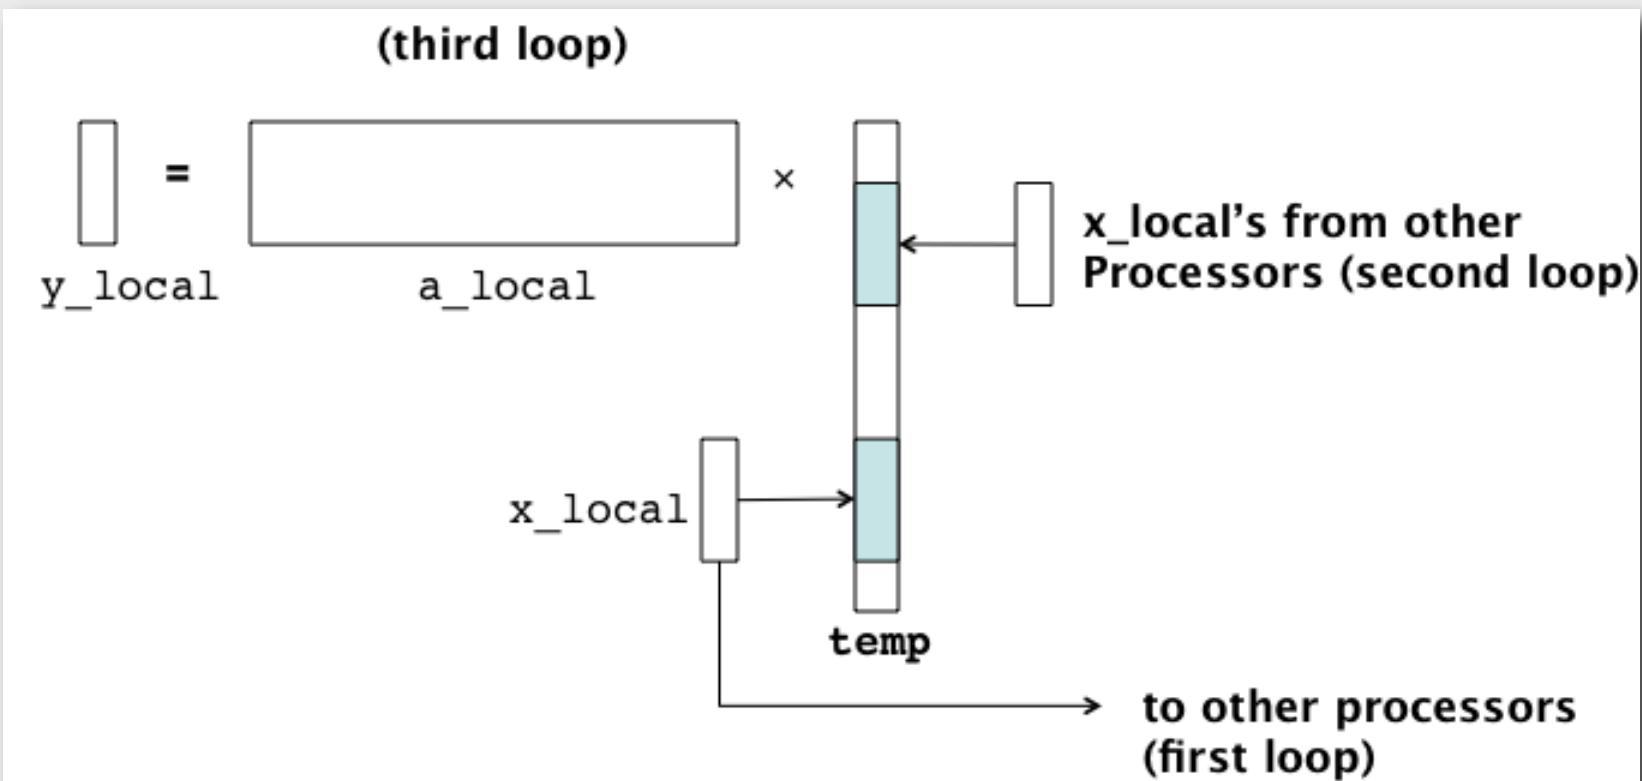
- Which processors gets which data?
  - A: chunk of rows,  $j = id * size .. (id + 1)*size - 1$
  - x: chunk of elements, same j
- Communication?
- Synchronization?

# Message Passing: MVP ver.01

Introduction to HPC  
in4049 TU Delft  
2013

```
float a_local[m/procs,n], y_local[m/procs], x_local[n/procs];
float temp[n];
int mypid;

mypid=proc(); k=n/procs;
for(i=0;i<procs;i++) {
    if(mypid==i) send(&x_local[0],k,i);
for(i=0;i<procs;i++) {
    if(mypid!=i)
        receive(&temp[i*k],k,i);
    else
        copy(temp[i*k],k,x_local)
}
for(i=0;i<k;i++)
    for(j=0;j<n;j++)
        y_local[i]=y_local[i]+a_local[i,j]*temp[j];
```



- Memory
  - *Need local buffer ( $\text{temp}[ ]$ ) of global size  $n$  anyway !*
- Performance issues?
  - *Communication with (proc-1) nodes*
    - *Sending  $n/\text{proc}$  items*
    - *Receiving  $n/\text{proc}$  items*
    - *Ordering of messages?*

# Message Passing: MPV – ver 02

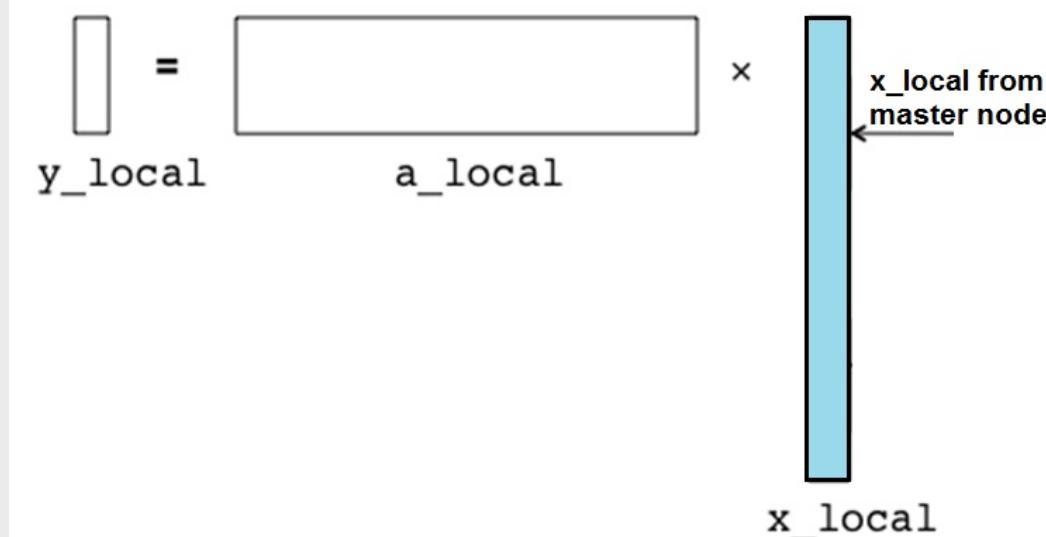
*Introduction to HPC  
in4049 TU Delft  
2013*

- another version of local structures results in:

```
float a_local[m/procs,n], y_local[m/procs],  
x_local[n];  
int mypid;  
  
mypid=proc(); k=n/procs;  
if (mypid==0)  
    send_all(&x_local[0],n);  
else  
    receive(&x_local[0],n,0);  
  
for(i=0;i<k;i++) {  
    for(j=0;j<n;j++) {  
        y_local[i]=y_local[i]+a_local[i,j]*x_local[j];  
    } }
```

# MPV: data structures - ver.02

Introduction to HPC  
in4049 TU Delft  
2013



- Memory
  - *No need for  $x\_local[n/procs]$*
  - *replace temp with full  $x\_local$*
- Performance issues?
  - *Broadcast : sending n items*
  - *Receiving n items*
  - *Ordering of messages?*

# Message Passing: Summary

Introduction to HPC  
in4049 TU Delft  
2013

- No notion of global data
- Data communication is done by message passing
  - Expensive, performance-wise
- Trade-off between:
  - One-copy data
    - *More communication is needed, less consistency issues*
  - Local data replication
    - *(apparently) Less communication, consistency is problematic*
- Techniques to improve performance:
  - Replicate read-only data
  - Computation and communication overlapping
  - Message aggregation

# Message Passing Quiz

Introduction to HPC  
in4049 TU Delft  
2013

Which of the following *conceptual programming models* be represented in message passing (MPI):

- *farmer/workers*
- *divide and conquer*
- *data parallelism*
- *task parallelism*
- *bulk synchronous*
- Write a short (pseudo-code) example for each one of those for which this is possible.
- Indicate the functionality
- Syntax correctness is not an issue.

---

Next lecture

---

## Next time – Lecture 4

---

*Introduction to HPC  
in4049 TU Delft  
2013*

- Parallel Programming Models (cont'd)
- Performance Analysis
  - Programming single-node
    - *Single CPU, Multi-core CPUs, GPUs*
    - Single node performance
    - Multi-node performance
  - Reading assignment (optional)
    - *M.D.Hill et. al - Amdhal's Law in the Multi-core Era (July '08)*
    - *V.W.Lee et. al - Debunking the 100x GPU vs CPU Myth: an evaluation of throughput computing on CPU and GPU (June '10)*